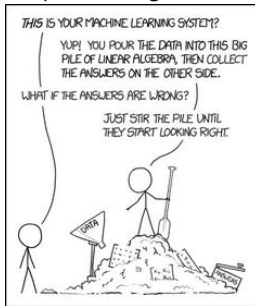


Deep Learning

Krystian Mikolajczyk & Carlo Ciliberto

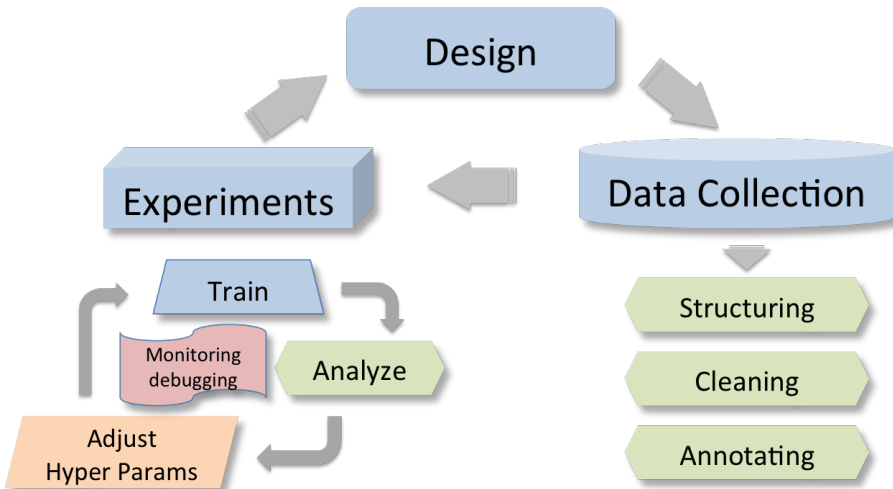
Department of Electrical and Electronic Engineering
Imperial College London



Practical development process

- System design/choice
- Data collection and augmentation
- Hyper parameters search
- Monitoring and debugging

Development process



Design

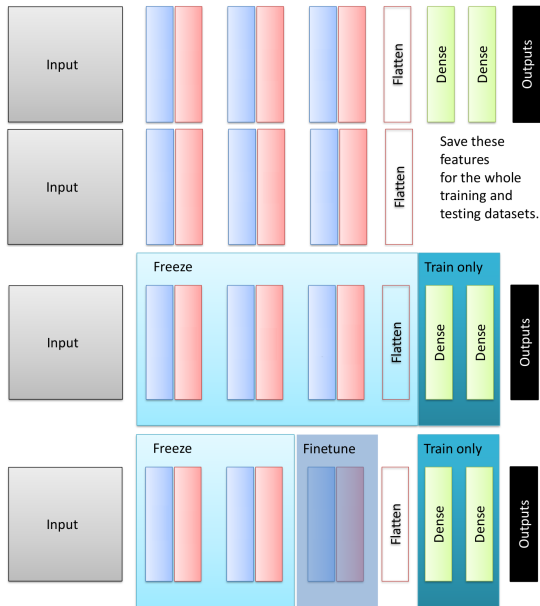
Baseline: end-to-end model

- Was the task studied before? Do literature review!
 - Start from competitions/survey papers
 - Establish a reasonable end-to-end system
- Choose the general category of model based on the structure of your data:
 - MLP for fixed size vectors
 - CNN for images
 - RNN for sequential data
- Nonlinear activations
 - Avoid sigmoid (except for output)
 - ReLU preferred (possibly Leaky ReLU)
 - Use Maxout if most ReLU units die (have zero activation)
- Weights & Biases
 - Random initialization with proper variance
 - For ReLU we prefer a small positive bias to activate ReLU

Design

Finetuning - borrow knowledge

- Pretrain your NN on a large dataset (e.g. same modality, similar task)
 - or start from a pretrained NN
- Option 1: remove / reshape the last few layers and use the features
- Option 2: Fine-tune the parameters on your own dataset
 - Freeze the parameters of first few layers, or make the learning rate small for them
 - Small data - train last FC layers only
 - Medium data - can finetune other layers
 - Use only 1/10th of the original learning rate in finetuning top layer, and 1/100th on intermediate layers



Data Collection

Study the application domain!

Collect data for the task

- How much data to collect?
 - The more the better
 - Depends on the effect we want to observe
 - Required error bounds and accuracy
- How to label the data?
 - Mechanical Turk, Freelancer, experts, ...
- Avoid bias
 - Selection, Sampling, ...

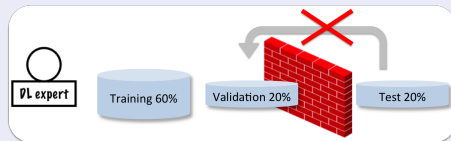
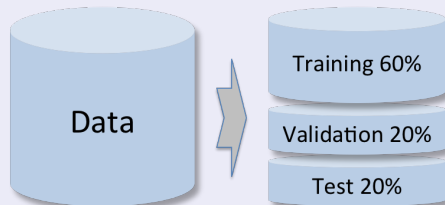
Dataset preparation/curation

- Data structuring & formatting
 - Is the data format suitable?
 - Standardization
- Data cleaning
 - incomplete data, anonymization, missing annotation, correction ...
- Data normalization
 - Value clipping/normalization
 - Whitening

Data Collection

Data split

- Training set
 - typically 60%, to run the learning algorithm on
 - Keep training data balanced
- Validation set
 - typically 20%, to tune hyper parameters, select features
 - make other decisions regarding the learning algorithm
 - also called development set
- Test set
 - typically 20%, to evaluate the performance of the algorithm



Testing

Do not use test data to make any decisions to improve learning!

Validation

Regularisation

$$\mathcal{L}_n(h) = \hat{R}_n(h) + \lambda \underbrace{\Omega(h)}$$

Validation

$$\underbrace{\mathcal{L}_n(h)} = \hat{R}_n(h) + \lambda \underbrace{\Omega(h)}$$

- 1 Split the training data \mathcal{D} into training \mathcal{D}_{train} and validation \mathcal{D}_{val} sets.
 overfit penalty *direct estimation* *overfit penalty*
- 2 Train g on \mathcal{D}_{train} .
- 3 Estimate its performance on \mathcal{D}_{val} ($v = |\mathcal{D}_{val}|$):

$$\check{R}_v(g) = \frac{1}{v} \sum_{(x_i, y_i) \in \mathcal{D}_{val}} \ell(g(x_i), y_i)$$

Very good estimate of $R(g)$

$$\mathbb{E}_{\mathcal{D}_{val}} [\check{R}_v(g)] \approx R(g) \leq \check{R}_v(g) + \underbrace{\Omega(v, \delta)}_{\sim \sqrt{\log(1/\delta)/v} \leftarrow \text{one model only on } v\text{-points}} \quad \text{w. p. } 1 - \delta$$

- \mathcal{D}_{val} is unbiased, small Hoeffding bound, only one g is considered.
- Select $\lambda^* = \operatorname{argmin}_{\lambda} \check{R}_v(g)$, then train on the whole \mathcal{D} with λ^* .

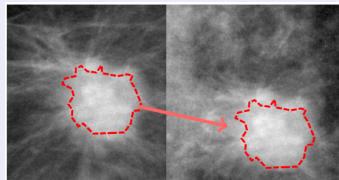
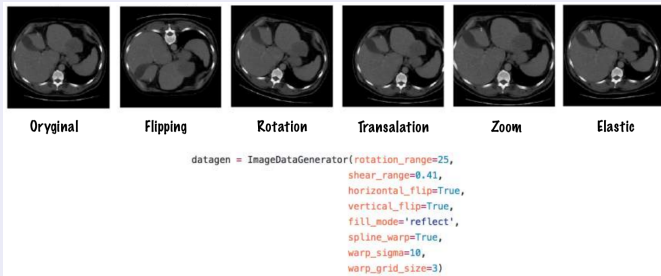
Data Augmentation

- Image data augmentation
- Text data augmentation
- Audio data augmentation

Data Augmentation

Image Data Augmentation

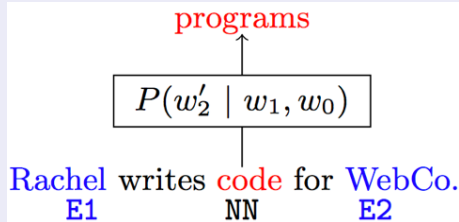
- Adding noise
- Generating modified samples
- Medical data
 - Segment tumor mass
 - Move
 - Resample background tissue
 - Blend



Data Augmentation

Text Data Augmentation

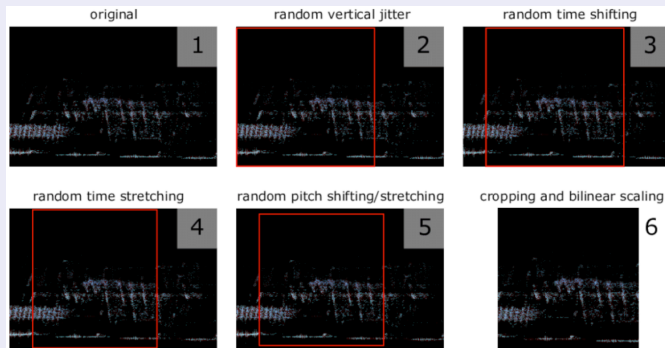
- Adding noise
- Inserting synonyms
- A conditional word-swap with externally trained language model and specifically targeting nouns (NN) between entity mentions (E1,E2)
- Rare words in new, synthetically created contexts



Data Augmentation

Audio Data Augmentation

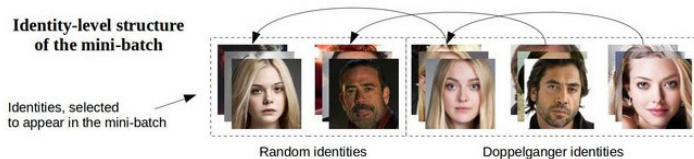
- Adding noise
- Vertical jitter
- Time shifting
- Time stretching: change the speed of the audio signal
- Pitch shift
- Cropping and bilinear scaling



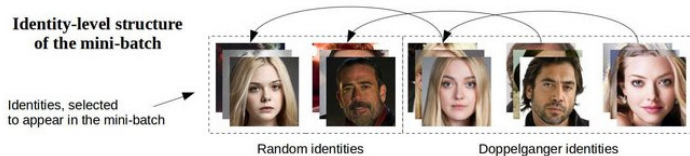
Hard Positive/Negative Mining

Adversarial training

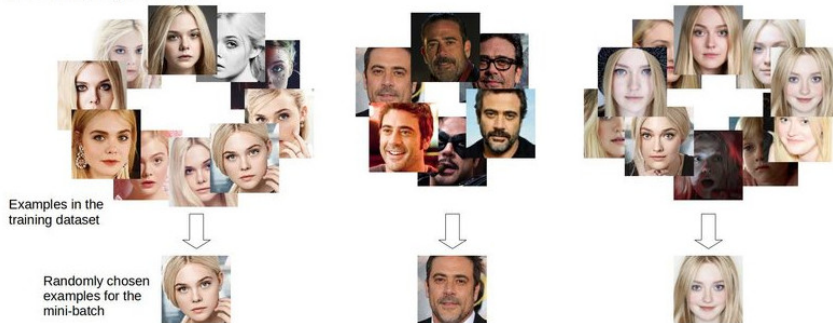
- Train the system from randomly formed mini-batches with balanced positive/negative examples
- Identify hard examples (close to decision boundary) during validation
- Form a new mini-batch by including the hard examples from the the previous iteration.



Hard Positive/Negative Mining

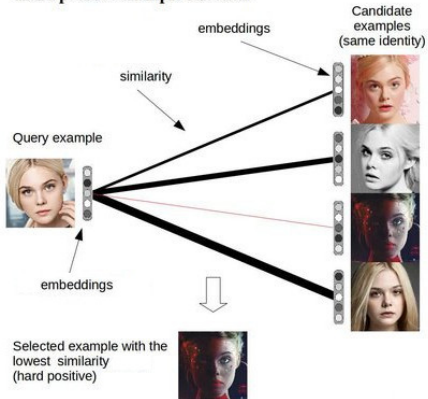


Random example selection

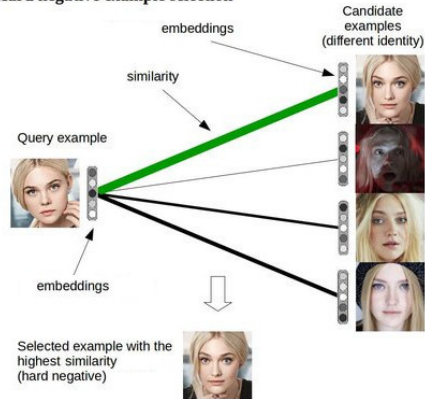


Hard Positive/Negative Mining


Hard positive example selection



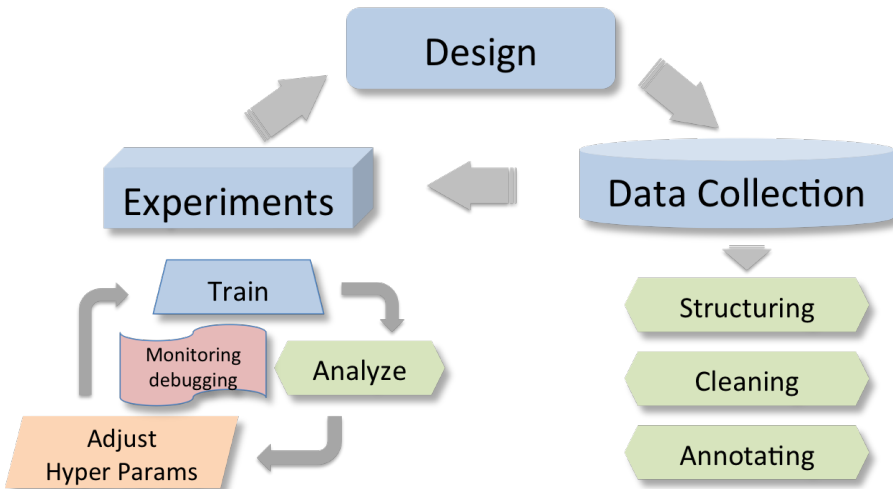
Hard negative example selection



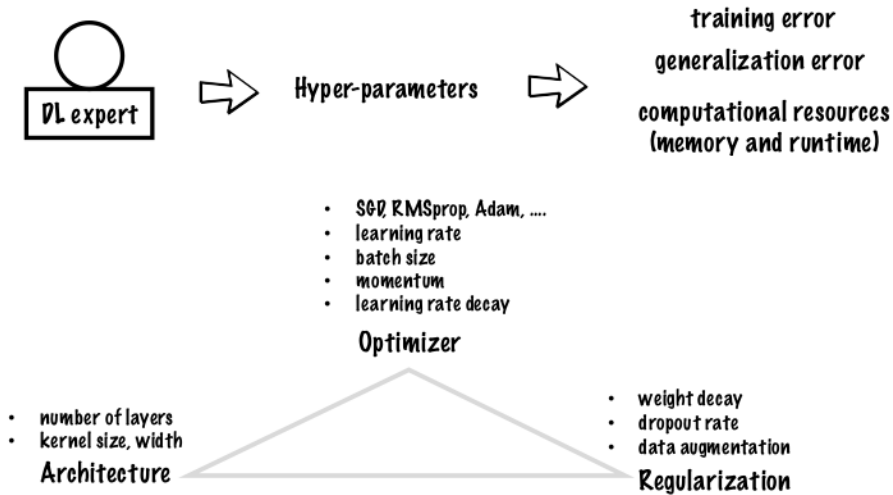
Example-level structure of the mini-batch



Development process



Parameter setting



Parameter setting

Hyper-parameters

- Optimization of hyper-parameters affects the quality of local minima that our model can reach (effective model capacity)
 - memory problems? reduce batch size
- The bigger the architecture the higher the model's capacity
 - caveat: memory and computation time are limited
- To decrease generalization gap increase regularization
 - increase weight decay and dropout rate to reduce model capacity
 - data augmentation does not affect model capacity

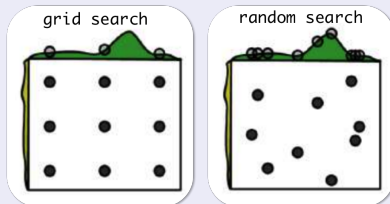
Good rule of thumb

Start from default parameters!

Parameter setting

Hyper parameter search: Good intuition and experience

- Learning Rate / Momentum
 - Decrease learning rate while training
 - Typical momentum to 0.8 - 0.9
- Batch Size
 - Shuffle data before batching
 - For large dataset: set to whatever fits your memory
 - For smaller dataset: find a tradeoff between instance randomness and gradient smoothness
- More efficient to optimize hyper-parameter with randomly chosen trials rather than on grid
- Model based hyper-parameter selection
- Use coarse to fine search for hyper-parameters
- Search on log scale (eg. learning rates)



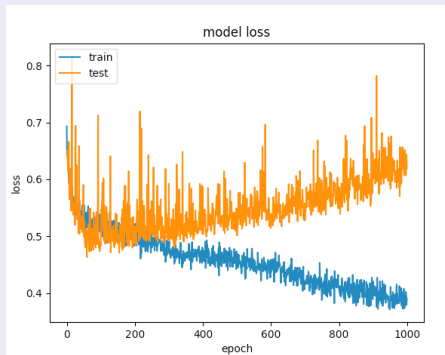
Monitoring and debugging

- Set a small dataset
- Reasoning using train and validation error/plots
- Monitoring histograms of activations and gradient updates
- Analysing of model predictions and errors

Monitoring and debugging

Set a small dataset

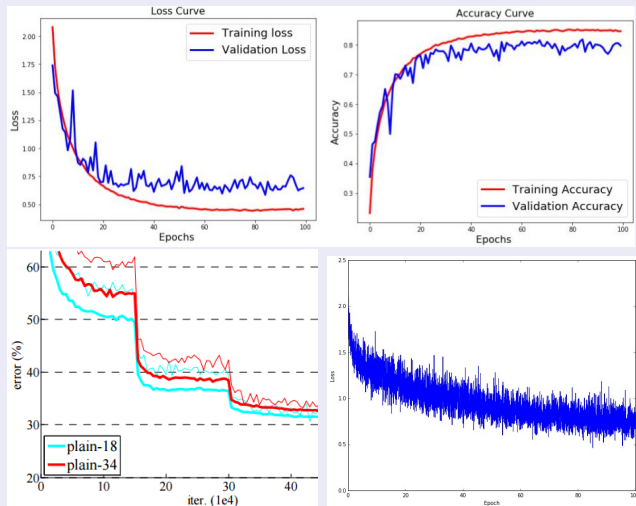
- Observe training error/plot
 - The model should overfit
- Change the optimisation parameters if the model does not overfit
 - If unsuccessful, inspect the code for potential bugs
- Reduce the size of the model (architecture) if model compilation time is long



Monitoring and debugging

- Optimize your hyper-parameter in validation and evaluate on test
- Keep track of training and validation loss during training
- Do early stopping if training and validation loss diverge
- Use patience - wait this long from last change
- Relative improvement threshold (significance)
- Loss does not tell you all. Try precision, class-wise precision, and other metrics

Train and validation error/plots



Monitoring and debugging

Visualize model predictions

- Check input data and annotations
- Check if predictions make sense

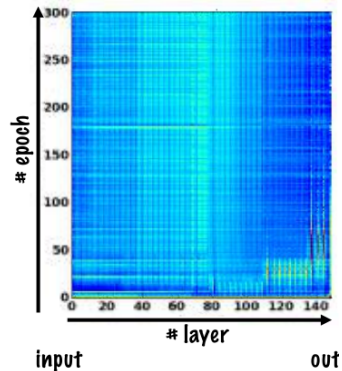
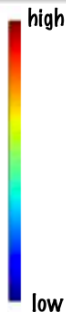
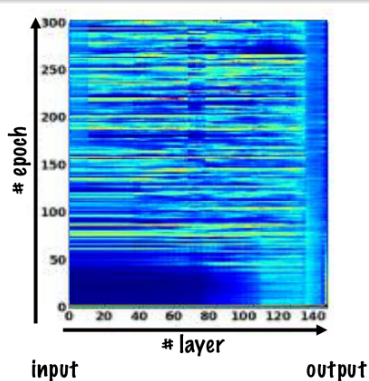
Understanding the underlying causes of the errors

- Get a sample of e.g. 100 validation set examples for which the system failed
- Examine these examples manually/visually
- Identify the most common errors
 - A simple excel table might be enough

Monitoring and debugging

Monitor histograms of activations and gradient updates per layer

- Erratic/unstable
- Uniform



Generalisation problem

- Biased data distributions (not i.i.d.)

Training & validation data



Test data



- Mix datasets and test
- Split validation into training validation and test validation

Training & validation data

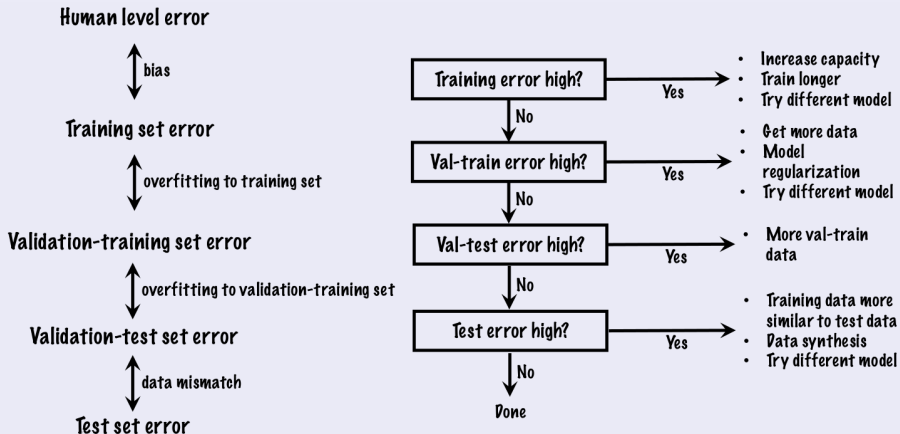


Test & validation data



Monitoring and debugging

Error analysis



Summary

- System design/choice
- Data collection and augmentation
- Hyper parameters search
- Monitoring and debugging